How Student Data Is Kept Off Claude's Servers

canvas-pii-guard - a plain-language security overview (proof of concept)

> **In one sentence:** when an instructor uses the Canvas content tool, only the **course material** they are building is ever sent to Claude. **Student rosters, grades, and submissions are stopped by software on the instructor's own computer** before any data could be sent.

The data path (what leaves the computer, and what doesn't)

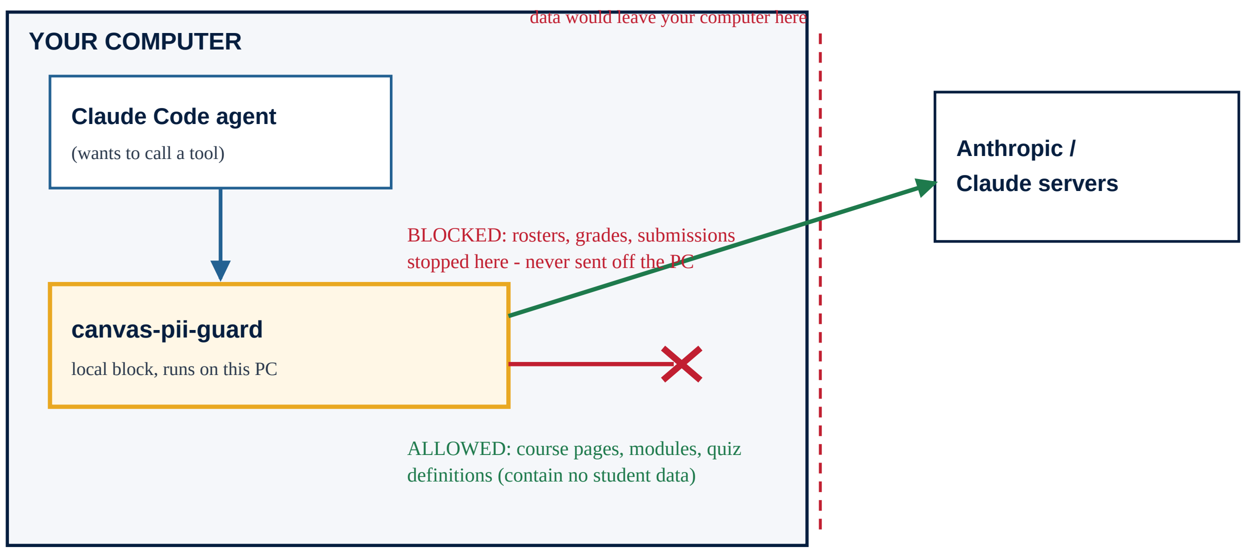

Every tool the agent runs passes through the local guard first. Course content continues to Claude; anything that looks like student data is blocked on the machine.

The layers of protection (defense in depth)

Layer	What it does	Type
1. AI policy	The assistant is instructed to refuse student data.	**Policy**
2. No risky tools	The content tool ships zero scripts that can read rosters/grades.	**Removes ability**
3. Local BLOCK hook	Software denies student-data requests before they run. **The guarantee.**	**Enforced**
4. Sterilizing gateway	If data is ever needed, raw stays on the PC; only scrubbed text is shown.	**Contained**
5. Output scrubber	Best-effort cleanup of stray IDs/emails in results.	**Backstop**

What happens in common situations

Situation	What the guard does	Result
Tool tries to pull a roster, grades, or submissions	Denied on the computer; no request is sent.	**BLOCKED**
Tool tries to open a saved student-data file	Denied.	**BLOCKED**

An unfamiliar student-data request	Denied by default (fail-safe).	**BLOCKED**
Building course pages / quizzes (no student data)	Allowed - this is the tool's job.	**ALLOWED**
A stray email or ID slips into output	Scrubbed automatically (best-effort).	**BACKSTOP**
Someone hides a request inside a custom script	Not caught by the guard alone.	**GAP***
Someone pastes a roster into the chat box	Not caught by the guard alone.	**GAP***

*Gaps are closed by the next step below (a Canvas login that simply cannot see grades). They are disclosed on purpose - see the contingency analysis.

Honest scope - please read

This is a proof of concept, and it is not a literal "air gap." The computer still uses the internet, and Claude still works normally. What this proves is that the **student-data path is blocked locally** for accidental and casual use - enforced by software you can read and test, not by trusting the AI. A determined person could still work around it. The single biggest hardening step (next phase) is to issue instructors a **Canvas login that cannot view grades or students at all**, which makes student data impossible to fetch at the source. Enforcing the block org-wide (so it can't be turned off) is the second step.

Why you can trust the claim (not just take our word)

• The rules are a short, readable file an IT reviewer can inspect.
• A built-in test runs **46 protection checks and passes all 46** (including MGCCC's real ID formats), and it openly lists what it does not catch.
• You can watch it work: with the guard on, a roster request is refused on the machine before any network call.